Date	Score	Passed
September 15, 2025 **11:43 PM IST**	110/110	Yes
Test_get_matrices	10/10	Show grader output
test_compute_loss	10/10	Show grader output
test_compute_gradient	10/10	Show grader output
test_align_embeddings	10/10	Show grader output
test_nearest_neighbor	10/10	Show grader output
unittest_test_vocabulary	10/10	Show grader output
test_get_document_embedding	10/10	Show grader output
test_get_document_vecs	10/10	Show grader output
test_hash_value_of_vector	10/10	Show grader output
test_make_hash_table	10/10	Show grader output
test_approximate_knn	10/10	Show grader output

Assignment 4 - Naive Machine Translation and LSH

You will now implement your first machine translation system and then you will see how locality sensitive hashing works. Let's get started by importing the required functions!

If you are running this notebook in your local computer, don't forget to download the twitter samples and stopwords from nltk.

```
nltk.download('stopwords')
nltk.download('twitter_samples')
```

Important Note on Submission to the AutoGrader

Before submitting your assignment to the AutoGrader, please make sure you are not doing the following:

1. You have not added any *extra* `print` statement(s) in the assignment.
2. You have not added any *extra* code cell(s) in the assignment.
3. You have not changed any of the function parameters.
4. You are not using any global variables inside your graded exercises. Unless specifically instructed to do so, please refrain from it and use the local variables instead.
5. You are not changing the assignment code where it is not required, like creating *extra* variables.

If you do any of the following, you will get something like, `Grader Error: Grader feedback not found` (or similarly unexpected) error upon submitting your assignment. Before asking for help/debugging the errors in your assignment, check for these first. If this is the case, and you don't remember the changes you have made, you can get a fresh copy of the assignment by following these [instructions (https://www.coursera.org/learn/classification-vector-spaces-in-nlp/supplement/YLuAg/h-ow-to-refresh-your-workspace)](https://www.coursera.org/learn/classification-vector-spaces-in-nlp/supplement/YLuAg/h-ow-to-refresh-your-workspace).

Table of Contents

In [5]:
```python
import pdb
import pickle
import string

import time

import nltk
import numpy as np
from nltk.corpus import stopwords, twitter_samples

from utils import (cosine_similarity, get_dict,
                   process_tweet)
from os import getcwd

import w4_unittest
```

In [6]:
```python
# add folder, tmp2, from our local workspace containing pre-downloaded corpor
filePath = f"{getcwd()}/tmp2/"
nltk.data.path.append(filePath)
```

1. The Word Embeddings Data for English and French Words

Write a program that translates English to French.

The Data

The full dataset for English embeddings is about 3.64 gigabytes, and the French embeddings are about 629 megabytes. To prevent the Coursera workspace from crashing, we've extracted a subset of the embeddings for the words that you'll use in this assignment.

The subset of data

To do the assignment on the Coursera workspace, we'll use the subset of word embeddings.

In [7]:
```python
en_embeddings_subset = pickle.load(open("./data/en_embeddings.p", "rb"))
fr_embeddings_subset = pickle.load(open("./data/fr_embeddings.p", "rb"))
```

Look at the data

- en_embeddings_subset: the key is an English word, and the value is a 300 dimensional array, which is the embedding for that word.

  ```
  'the': array([ 0.08007812,  0.10498047,  0.04980469,  0.0534668 ,
         -0.06738281, ....
  ```

- fr_embeddings_subset: the key is a French word, and the value is a 300 dimensional array, which is the embedding for that word.

  ```
  'la': array([-6.18250e-03, -9.43867e-04, -8.82648e-03,  3.24623e-
         03
  ```

Load two dictionaries mapping the English to French words

- A training dictionary
- and a testing dictionary.

In [8]:
```
# loading the english to french dictionaries
en_fr_train = get_dict('./data/en-fr.train.txt')
print('The length of the English to French training dictionary is', len(en_fr_
en_fr_test = get_dict('./data/en-fr.test.txt')
print('The length of the English to French test dictionary is', len(en_fr_test
```

```
The length of the English to French training dictionary is 5000
The length of the English to French test dictionary is 1500
```

Looking at the English French dictionary

- `en_fr_train` is a dictionary where the key is the English word and the value is the French translation of that English word.

  ```
  {'the': 'la',
   'and': 'et',
   'was': 'était',
   'for': 'pour',
  ```

- `en_fr_test` is similar to `en_fr_train`, but is a test set. We won't look at it until we get to testing.

1.1 Generate Embedding and Transform Matrices

Exercise 1 - get_matrices

Translating English dictionary to French by using embeddings.

You will now implement a function `get_matrices`, which takes the loaded data and returns matrices `X` and `Y`.

Inputs:

- `en_fr` : English to French dictionary
- `en_embeddings` : English to embeddings dictionary

- `fr_embeddings` : French to embeddings dictionary

Returns:

- Matrix `X` and matrix `Y`, where each row in X is the word embedding for an english word, and the same row in Y is the word embedding for the French version of that English word.

$$X$$

English word embedding $\left\{\begin{array}{cccccc} 2 & 9 & \cdots & \cdots & 1 & 3 \\ 1 & 2 & \cdots & \cdots & 0 & 3 \\ 9 & 8 & \cdots & \cdots & 7 & 5 \\ 1 & 0 & \cdots & \cdots & 9 & 3 \\ 4 & 9 & \cdots & \cdots & 4 & 0 \\ 2 & 1 & \cdots & \cdots & 0 & 5 \end{array}\right.$

$$Y$$

French word embedding $\left\{\begin{array}{cccccc} 1 & 9 & \cdots & \cdots & 9 & 3 \\ 4 & 2 & \cdots & \cdots & 0 & 8 \\ 9 & 6 & \cdots & \cdots & 6 & 5 \\ 1 & 9 & \cdots & \cdots & 9 & 3 \\ 4 & 2 & \cdots & \cdots & 4 & 0 \\ 2 & 8 & \cdots & \cdots & 6 & 5 \end{array}\right.$
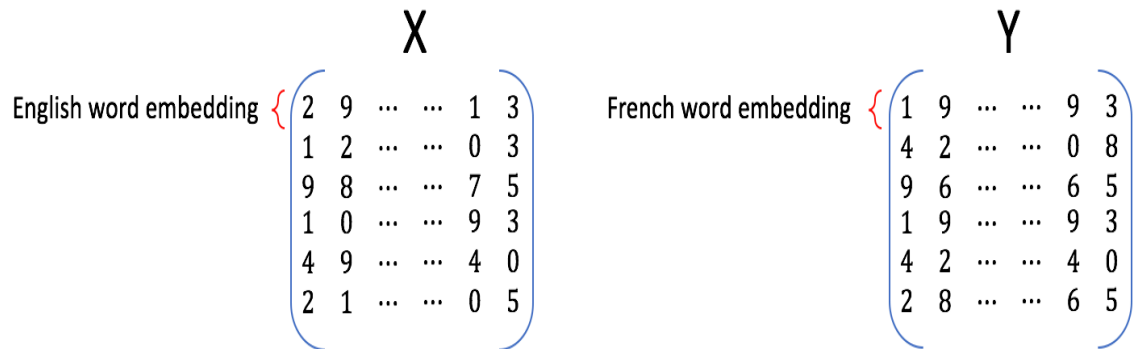
Figure 1

Use the `en_fr` dictionary to ensure that the ith row in the `X` matrix corresponds to the ith row in the `Y` matrix.

Instructions: Complete the function `get_matrices()` :

- Iterate over English words in `en_fr` dictionary.
- Check if the word have both English and French embedding.

Hints

```python
# UNQ_C1 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def get_matrices(en_fr, french_vecs, english_vecs):
    """
    Creates matrices of word embeddings for English and French words that are

    Inputs:
        en_fr: Dictionary mapping English words to French words.
        french_vecs: Dictionary of French word embeddings.
        english_vecs: Dictionary of English word embeddings.

    Outputs:
        X: Matrix with each row being the embedding of an English word. Shape
        Y: Matrix with each row being the embedding of the corresponding Frenc

    Note:
        This function does not compute or return a projection matrix.
    """

    ### START CODE HERE ###

    # X_l and Y_l are lists of the english and french word embeddings
    X_l = list()
    Y_l = list()

    # get the english words (the keys in the dictionary) and store in a set()
    english_set = None

    # get the french words (keys in the dictionary) and store in a set()
    french_set = None

    # store the french words that are part of the english-french dictionary (t
    french_words = set(en_fr.values())

    # loop through all english, french word pairs in the english french dictic
    for en_word, fr_word in en_fr.items():

        # check that the french word has an embedding and that the english wor
        if fr_word in french_set and en_word in english_set:

            # get the english embedding
            en_vec = english_vecs[en_word]

            # get the french embedding
            fr_vec = None

            # add the english embedding to the list
            X_l.append(en_vec)

            # add the french embedding to the list
            None

    # stack the vectors of X_l into a matrix X
    X = None

    # stack the vectors of Y_l into a matrix Y
    Y = None
    ### END CODE HERE ###
```

```
        return X, Y
```

Now we will use function `get_matrices()` to obtain sets `X_train` and `Y_train` of English and French word embeddings into the corresponding vector space models.

```
In [ ]:   # UNQ_C2 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
          # You do not have to input any code in this cell, but it is relevant to gradin

          # getting the training set:
          X_train, Y_train = get_matrices(
              en_fr_train, fr_embeddings_subset, en_embeddings_subset)
```

```
In [ ]:   # Test your function
          w4_unittest.test_get_matrices(get_matrices)
```

2 - Translations

English   Natural Language Processing ⟶   Francais

Figure 2

Write a program that translates English words to French words using word embeddings and vector space models.

2.1 - Translation as Linear Transformation of Embeddings

Given dictionaries of English and French word embeddings you will create a transformation matrix `R`

- Given an English word embedding, $\mathbf{e}$, you can multiply $\mathbf{eR}$ to get a new word embedding $\mathbf{f}$.
    - Both $\mathbf{e}$ and $\mathbf{f}$ are [row vectors (https://en.wikipedia.org/wiki/Row_and_column_vectors)](https://en.wikipedia.org/wiki/Row_and_column_vectors).
- You can then compute the nearest neighbors to `f` in the french embeddings and recommend the word that is most similar to the transformed word embedding.

Describing translation as the minimization problem

Find a matrix `R` that minimizes the following equation.

$$\arg \min_{\mathbf{R}} \|\mathbf{XR} - \mathbf{Y}\|_F \tag{1}$$

Frobenius norm

The Frobenius norm of a matrix A (assuming it is of dimension m, n) is defined as the square root of the sum of the absolute squares of its elements:

$$\|\mathbf{A}\|_F \equiv \sqrt{\sum_{i=1}^{m} \sum_{j=1}^{n} |a_{ij}|^2} \tag{2}$$

Actual loss function

In the real world applications, the Frobenius norm loss:

$$\|\mathbf{XR} - \mathbf{Y}\|_F$$

is often replaced by it's squared value divided by m:

$$\frac{1}{m}\|\mathbf{XR} - \mathbf{Y}\|_F^2$$

where m is the number of examples (rows in $\mathbf{X}$).

- The same R is found when using this loss function versus the original Frobenius norm.
- The reason for taking the square is that it's easier to compute the gradient of the squared Frobenius.
- The reason for dividing by m is that we're more interested in the average loss per embedding than the loss for the entire training set.
  - The loss for all training set increases with more words (training examples), so taking the average helps us to track the average loss regardless of the size of the training set.

[Optional] Detailed explanation why we use norm squared instead of the norm:
Click for optional details

Implementing translation mechanism described in this section.

Exercise 2 - compute_loss

Step 1: Computing the loss

- The loss function will be squared Frobenius norm of the difference between matrix and its approximation, divided by the number of training examples m.
- Its formula is:

$$L(X, Y, R) = \frac{1}{m} \sum_{i=1}^{m} \sum_{j=1}^{n} \left(a_{ij}\right)^2$$

where a_{ij} is value in ith row and jth column of the matrix $\mathbf{XR} - \mathbf{Y}$.

Instructions: complete the `compute_loss()` function

- Compute the approximation of `Y` by matrix multiplying `X` and `R`

- Compute difference `XR - Y`
- Compute the squared Frobenius norm of the difference and divide it by m.

Hints

```
In [12]:  import numpy as np

          def compute_loss(X, Y, R):
              '''
              Inputs:
                  X: a matrix of dimension (m,n) where the columns are the English embed
                  Y: a matrix of dimension (m,n) where the columns correspong to the Fre
                  R: a matrix of dimension (n,n) - transformation matrix from English to
              Outputs:
                  L: a scalar - the value of the loss function for given X, Y and R.
              '''
              ### START CODE HERE ###
              m = X.shape[0]                  # number of rows / examples
              diff = X.dot(R) - Y          # XR - Y
              diff_squared = diff**2       # element-wise square
              sum_diff_squared = np.sum(diff_squared)  # sum all elements
              loss = sum_diff_squared / m                # average over examples
              ### END CODE HERE ###
              return loss
```

```
In [13]:  # Testing your implementation.
          np.random.seed(123)
          m = 10
          n = 5
          X = np.random.rand(m, n)
          Y = np.random.rand(m, n) * .1
          R = np.random.rand(n, n)
          print(f"Expected loss for an experiment with random matrices: {compute_loss(X,
```

Expected loss for an experiment with random matrices: 8.1866

Expected output:

```
Expected loss for an experiment with random matrices: 8.1866
```

```
In [14]:  # Test your function
          w4_unittest.test_compute_loss(compute_loss)
```

All tests passed

Exercise 3 - compute_gradient

Step 2: Computing the gradient of loss with respect to transform matrix R

- Calculate the gradient of the loss with respect to transform matrix `R`.

- The gradient is a matrix that encodes how much a small change in R affect the change in the loss function.
- The gradient gives us the direction in which we should decrease R to minimize the loss.
- m is the number of training examples (number of rows in X).
- The formula for the gradient of the loss function $L(X, Y, R)$ is:

$$\frac{d}{dR}L(X, Y, R) = \frac{d}{dR}\left(\frac{1}{m}\|XR - Y\|_F^2\right) = \frac{2}{m}X^T(XR - Y)$$

Hints

```
In [15]:  # UNQ_C4
          import numpy as np

          def compute_gradient(X, Y, R):
              '''
              Inputs:
                  X: a matrix of dimension (m,n) where the columns are the English embed
                  Y: a matrix of dimension (m,n) where the columns correspond to the Fre
                  R: a matrix of dimension (n,n) - transformation matrix from English to
              Outputs:
                  g: a matrix of dimension (n,n) - gradient of the loss function L for g
              '''
              ### START CODE HERE ###
              m = X.shape[0]                      # number of examples
              gradient = (2/m) * X.T.dot(X.dot(R) - Y)
              ### END CODE HERE ###
              return gradient
```

```
In [16]:  # Testing your implementation.
          np.random.seed(123)
          m = 10
          n = 5
          X = np.random.rand(m, n)
          Y = np.random.rand(m, n) * .1
          R = np.random.rand(n, n)
          gradient = compute_gradient(X, Y, R)
          print(f"First row of the gradient matrix: {gradient[0]}")
```

```
First row of the gradient matrix: [1.3498175  1.11264981 0.69626762 0.984684
99 1.33828969]
```

Expected output:

```
First row of the gradient matrix: [1.3498175  1.11264981 0.69626762
0.98468499 1.33828969]
```

```
In [17]:  # Test your function
          w4_unittest.test_compute_gradient(compute_gradient)
```

```
All tests passed
```

Step 3: Finding the optimal R with Gradient Descent Algorithm

Gradient Descent

[Gradient descent (https://ml-cheatsheet.readthedocs.io/en/latest/gradient_descent.html)](https://ml-cheatsheet.readthedocs.io/en/latest/gradient_descent.html) is an iterative algorithm which is used in searching for the optimum of the function.

- Earlier, we've mentioned that the gradient of the loss with respect to the matrix encodes how much a tiny change in some coordinate of that matrix affect the change of loss function.
- Gradient descent uses that information to iteratively change matrix `R` until we reach a point where the loss is minimized.

Training with a fixed number of iterations

Most of the time we iterate for a fixed number of training steps rather than iterating until the loss falls below a threshold.

OPTIONAL: explanation for fixed number of iterations
click here for detailed discussion

Pseudocode:

1. Calculate gradient g of the loss with respect to the matrix R.
2. Update R with the formula:
$$R_{\text{new}} = R_{\text{old}} - \alpha g$$

Where α is the learning rate, which is a scalar.

Learning Rate

- The learning rate or "step size" α is a coefficient which decides how much we want to change R in each step.
- If we change R too much, we could skip the optimum by taking too large of a step.
- If we make only small changes to R, we will need many steps to reach the optimum.
- Learning rate α is used to control those changes.
- Values of α are chosen depending on the problem, and we'll use `learning_rate` $= 0.0003$ as the default value for our algorithm.

Exercise 4 - align_embeddings

Implement `align_embeddings()`

Hints

```python
In [18]:  # UNQ_C5
          import numpy as np

          def align_embeddings(X, Y, train_steps=100, learning_rate=0.0003, verbose=True
              '''
              Inputs:
                  X: a matrix of dimension (m,n) where the columns are the English embed
                  Y: a matrix of dimension (m,n) where the columns correspond to the Fre
                  train_steps: number of gradient descent steps.
                  learning_rate: step size for gradient descent.
              Outputs:
                  R: a matrix of dimension (n,n) - projection matrix minimizing ||XR - Y
              '''
              np.random.seed(129)
              R = np.random.rand(X.shape[1], X.shape[1])

              for i in range(train_steps):
                  if verbose and i % 25 == 0:
                      print(f"loss at iteration {i} is: {compute_loss(X, Y, R):.4f}")

                      ### START CODE HERE ###
                      gradient = compute_gradient(X, Y, R)
                      R -= learning_rate * gradient
                      ### END CODE HERE ###

              return R
```

```python
In [19]:  # UNQ_C6 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
          # You do not have to input any code in this cell, but it is relevant to gradin

          # Testing your implementation.
          np.random.seed(129)
          m = 10
          n = 5
          X = np.random.rand(m, n)
          Y = np.random.rand(m, n) * .1
          R = align_embeddings(X, Y)
```

```
loss at iteration 0 is: 3.7242
loss at iteration 25 is: 3.6283
loss at iteration 50 is: 3.5350
loss at iteration 75 is: 3.4442
```

Expected Output:

```
loss at iteration 0 is: 3.7242
loss at iteration 25 is: 3.6283
loss at iteration 50 is: 3.5350
loss at iteration 75 is: 3.4442
```

```python
In [20]:  # Test your function
          w4_unittest.test_align_embeddings(align_embeddings)
```

```
All tests passed
```

Calculate Transformation matrix R

Using just the training set, find the transformation matrix R by calling the function `align_embeddings()`.

NOTE: The code cell below will take a few minutes to fully execute (~3 mins)

```
In [21]: # UNQ_C7 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
         # You do not have to input any code in this cell, but it is relevant to gradin
         R_train = align_embeddings(X_train, Y_train, train_steps=400, learning_rate=0.
```

```
loss at iteration 0 is: 963.0146
loss at iteration 25 is: 97.8292
loss at iteration 50 is: 26.8329
loss at iteration 75 is: 9.7893
loss at iteration 100 is: 4.3776
loss at iteration 125 is: 2.3281
loss at iteration 150 is: 1.4480
loss at iteration 175 is: 1.0338
loss at iteration 200 is: 0.8251
loss at iteration 225 is: 0.7145
loss at iteration 250 is: 0.6534
loss at iteration 275 is: 0.6185
loss at iteration 300 is: 0.5981
loss at iteration 325 is: 0.5858
loss at iteration 350 is: 0.5782
loss at iteration 375 is: 0.5735
```

Expected Output

```
loss at iteration 0 is: 963.0146
loss at iteration 25 is: 97.8292
loss at iteration 50 is: 26.8329
loss at iteration 75 is: 9.7893
loss at iteration 100 is: 4.3776
loss at iteration 125 is: 2.3281
loss at iteration 150 is: 1.4480
loss at iteration 175 is: 1.0338
loss at iteration 200 is: 0.8251
loss at iteration 225 is: 0.7145
loss at iteration 250 is: 0.6534
loss at iteration 275 is: 0.6185
loss at iteration 300 is: 0.5981
loss at iteration 325 is: 0.5858
loss at iteration 350 is: 0.5782
loss at iteration 375 is: 0.5735
```

2.2 - Testing the Translation

k-Nearest Neighbors Algorithm

- k-NN is a method which takes a vector as input and finds the other vectors in the dataset that are closest to it.
- The 'k' is the number of "nearest neighbors" to find (e.g. k=2 finds the closest two neighbors).

Searching for the Translation Embedding

Since we're approximating the translation function from English to French embeddings by a linear transformation matrix $\mathbf{R}$, most of the time we won't get the exact embedding of a French word when we transform embedding $\mathbf{e}$ of some particular English word into the French embedding space.

- This is where k-NN becomes really useful! By using 1-NN with $\mathbf{eR}$ as input, we can search for an embedding $\mathbf{f}$ (as a row) in the matrix $\mathbf{Y}$ which is the closest to the transformed vector $\mathbf{eR}$

Cosine Similarity

Cosine similarity between vectors u and v calculated as the cosine of the angle between them. The formula is

$$\cos(u, v) = \frac{u \cdot v}{\|u\| \, \|v\|}$$

- $\cos(u, v) = 1$ when u and v lie on the same line and have the same direction.
- $\cos(u, v)$ is -1 when they have exactly opposite directions.
- $\cos(u, v)$ is 0 when the vectors are orthogonal (perpendicular) to each other.

Note: Distance and similarity are pretty much opposite things.

- We can obtain distance metric from cosine similarity, but the cosine similarity can't be used directly as the distance metric.
- When the cosine similarity increases (towards 1), the "distance" between the two vectors decreases (towards 0).
- We can define the cosine distance between u and v as
$$d_{\cos}(u, v) = 1 - \cos(u, v)$$

Exercise 5 - nearest_neighbor

Complete the function `nearest_neighbor()`

Inputs:

- Vector `v` ,
- A set of possible nearest neighbors `candidates`
- `k` nearest neighbors to find.
- The distance metric should be based on cosine similarity.
- `cosine_similarity` function is already implemented and imported for you. It's arguments are two vectors and it returns the cosine of the angle between them.

- Iterate over rows in `candidates`, and save the result of similarities between current row and vector `v` in a python list. Take care that similarities are in the same order as row vectors of `candidates`.
- Now you can use [numpy argsort (https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html#numpy.argsort)](https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html#numpy.argsort) to sort the indices for the rows of `candidates`.

Hints

```python
In [22]:  # UNQ_C8
          import numpy as np

          def nearest_neighbor(v, candidates, k=1, cosine_similarity=cosine_similarity):
              """
              Input:
                - v: the vector you want to find the nearest neighbor for
                - candidates: a set of vectors where we will find the neighbors
                - k: top k nearest neighbors to find
              Output:
                - k_idx: the indices of the top k closest vectors in sorted form
              """
              ### START CODE HERE ###
              similarity_l = []

              for row in candidates:
                  cos_similarity = cosine_similarity(v, row)
                  similarity_l.append(cos_similarity)

              similarity_l = np.array(similarity_l)
              sorted_ids = np.argsort(similarity_l)   # ascending order
              sorted_ids = sorted_ids[::-1]           # descending order
              k_idx = sorted_ids[:k]
              ### END CODE HERE ###
              return k_idx
```

```python
In [23]:  # UNQ_C9 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
          # You do not have to input any code in this cell, but it is relevant to gradin

          # Test your implementation:
          v = np.array([1, 0, 1])
          candidates = np.array([[1, 0, 5], [-2, 5, 3], [2, 0, 1], [6, -9, 5], [9, 9, 9]
          print(candidates[nearest_neighbor(v, candidates, 3)])
```

```
[[2 0 1]
 [1 0 5]
 [9 9 9]]
```

Expected Output:

```
[[2 0 1]
 [1 0 5]
 [9 9 9]]
```

```
In [24]:  # Test your function
          w4_unittest.test_nearest_neighbor(nearest_neighbor)
```

All tests passed

Test your Translation and Compute its Accuracy

Exercise 6 - test_vocabulary

Complete the function `test_vocabulary` which takes in English embedding matrix X, French embedding matrix Y and the R matrix and returns the accuracy of translations from X to Y by R.

- Iterate over transformed English word embeddings and check if the closest French word vector belongs to French word that is the actual translation.
- Obtain an index of the closest French embedding by using `nearest_neighbor` (with argument `k=1`), and compare it to the index of the English embedding you have just transformed.
- Keep track of the number of times you get the correct translation.
- Calculate accuracy as

$$\text{accuracy} = \frac{\#(\text{correct predictions})}{\#(\text{total predictions})}$$

```
In [25]:  # UNQ_C10
          def test_vocabulary(X, Y, R, nearest_neighbor=nearest_neighbor):
              '''
              Input:
                  X: a matrix where the rows are the English embeddings.
                  Y: a matrix where the rows correspond to the French embeddings.
                  R: the transform matrix mapping English embeddings to French embedding
              Output:
                  accuracy: proportion of correct nearest neighbor matches for English→F
              '''
              ### START CODE HERE ###
              pred = X.dot(R)   # transform English embeddings

              num_correct = 0

              for i in range(len(pred)):
                  pred_idx = nearest_neighbor(pred[i], Y, k=1)[0]   # get index of neares
                  if pred_idx == i:
                      num_correct += 1

              accuracy = num_correct / len(pred)
              ### END CODE HERE ###

              return accuracy
```

Let's see how is your translation mechanism working on the unseen data:

```
In [26]:  X_val, Y_val = get_matrices(en_fr_test, fr_embeddings_subset, en_embeddings_su
```

```
In [27]:  # UNQ_C11 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
          # You do not have to input any code in this cell, but it is relevant to gradin

          acc = test_vocabulary(X_val, Y_val, R_train)  # this might take a minute or tw
          print(f"accuracy on test set is {acc:.3f}")
```

accuracy on test set is 0.557

Expected Output:

```
0.557
```

You managed to translate words from one language to another language without ever seeing them with almost 56% accuracy by using some basic linear algebra and learning a mapping of words from one language to another!

```
In [28]:  # Test your function
          w4_unittest.unittest_test_vocabulary(test_vocabulary)
```

All tests passed

3 - LSH and Document Search

In this part of the assignment, you will implement a more efficient version of k-nearest neighbors using locality sensitive hashing. You will then apply this to document search.

- Process the tweets and represent each tweet as a vector (represent a document with a vector embedding).
- Use locality sensitive hashing and k nearest neighbors to find tweets that are similar to a given tweet.

```
In [29]:  # get the positive and negative tweets
          all_positive_tweets = twitter_samples.strings('positive_tweets.json')
          all_negative_tweets = twitter_samples.strings('negative_tweets.json')
          all_tweets = all_positive_tweets + all_negative_tweets
```

3.1 - Getting the Document Embeddings

Bag-of-words (BOW) Document Models

Text documents are sequences of words.

- The ordering of words makes a difference. For example, sentences "Apple pie is better than pepperoni pizza." and "Pepperoni pizza is better than apple pie" have opposite meanings due to the word ordering.
- However, for some applications, ignoring the order of words can allow us to train an efficient and still effective model.

- This approach is called Bag-of-words document model.

Document Embeddings

- Document embedding is created by summing up the embeddings of all words in the document.
- If we don't know the embedding of some word, we can ignore that word.

Exercise 7 - get_document_embedding

Complete the `get_document_embedding()` function.

- The function `get_document_embedding()` encodes entire document as a "document" embedding.
- It takes in a document (as a string) and a dictionary, `en_embeddings`
- It processes the document, and looks up the corresponding embedding of each word.
- It then sums them up and returns the sum of all word vectors of that processed tweet.

Hints

```python
In [30]:   # UNQ_C12
           import numpy as np

           def get_document_embedding(tweet, en_embeddings, process_tweet=process_tweet):
               '''
               Input:
                   - tweet: a string
                   - en_embeddings: a dictionary of word embeddings
               Output:
                   - doc_embedding: sum of all word embeddings in the tweet
               '''
               doc_embedding = np.zeros(300)

               ### START CODE HERE ###
               processed_doc = process_tweet(tweet)

               for word in processed_doc:
                   if word in en_embeddings:
                       doc_embedding += en_embeddings[word]
               ### END CODE HERE ###

               return doc_embedding
```

```python
In [31]:   # UNQ_C13 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
           # You do not have to input any code in this cell, but it is relevant to gradin

           # testing your function
           custom_tweet = "RT @Twitter @chapagain Hello There! Have a great day. :) #good
           tweet_embedding = get_document_embedding(custom_tweet, en_embeddings_subset)
           tweet_embedding[-5:]
```

```
Out[31]:   array([-0.00268555, -0.15378189, -0.55761719, -0.07216644, -0.32263184])
```

Expected output:

```
array([-0.00268555, -0.15378189, -0.55761719, -0.07216644, -0.3226318
4])
```

In [32]: `# Test your function`
`w4_unittest.test_get_document_embedding(get_document_embedding)`

All tests passed

Exercise 8 - get_document_vecs

Store all document vectors into a dictionary

Now, let's store all the tweet embeddings into a dictionary. Implement `get_document_vecs()`

In [33]:
```python
# UNQ_C14
import numpy as np

def get_document_vecs(all_docs, en_embeddings, get_document_embedding=get_docu
    '''
    Input:
        - all_docs: list of strings - all tweets in our dataset.
        - en_embeddings: dictionary with words as the keys and their embedding
    Output:
        - document_vec_matrix: matrix of tweet embeddings.
        - ind2Doc_dict: dictionary with indices of tweets in vecs as keys and
    '''

    ind2Doc_dict = {}
    document_vec_l = []

    for i, doc in enumerate(all_docs):
        ### START CODE HERE ###
        doc_embedding = get_document_embedding(doc, en_embeddings)
        ind2Doc_dict[i] = doc_embedding
        document_vec_l.append(doc_embedding)
        ### END CODE HERE ###

    document_vec_matrix = np.vstack(document_vec_l)

    return document_vec_matrix, ind2Doc_dict
```

In [34]: `document_vecs, ind2Tweet = get_document_vecs(all_tweets, en_embeddings_subset)`

```
In [35]:   # UNQ_C15 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
           # You do not have to input any code in this cell, but it is relevant to gradin

           print(f"length of dictionary {len(ind2Tweet)}")
           print(f"shape of document_vecs {document_vecs.shape}")
```

```
length of dictionary 10000
shape of document_vecs (10000, 300)
```

Expected Output

```
length of dictionary 10000
shape of document_vecs (10000, 300)
```

```
In [36]:   # Test your function. This cell may take some seconds to run.
           w4_unittest.test_get_document_vecs(get_document_vecs)
```

```
All tests passed
```

3.2 - Looking up the Tweets

Now you have a vector of dimension (m,d) where m is the number of tweets (10,000) and d
is the dimension of the embeddings (300). Now you will input a tweet, and use cosine similarity
to see which tweet in our corpus is similar to your tweet.

```
In [37]:   my_tweet = 'i am sad'
           process_tweet(my_tweet)
           tweet_embedding = get_document_embedding(my_tweet, en_embeddings_subset)
```

```
In [38]:   # UNQ_C16 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
           # You do not have to input any code in this cell, but it is relevant to gradin

           # this gives you a similar tweet as your input.
           # this implementation is vectorized...
           idx = np.argmax(cosine_similarity(document_vecs, tweet_embedding))
           print(all_tweets[idx])
```

```
@hanbined sad pray for me :(((
```

Expected Output

```
@hanbined sad pray for me :(((
```

3.3 - Finding the most Similar Tweets with LSH

You will now implement locality sensitive hashing (LSH) to identify the most similar tweet.

- Instead of looking at all 10,000 vectors, you can just search a subset to find its nearest neighbors.

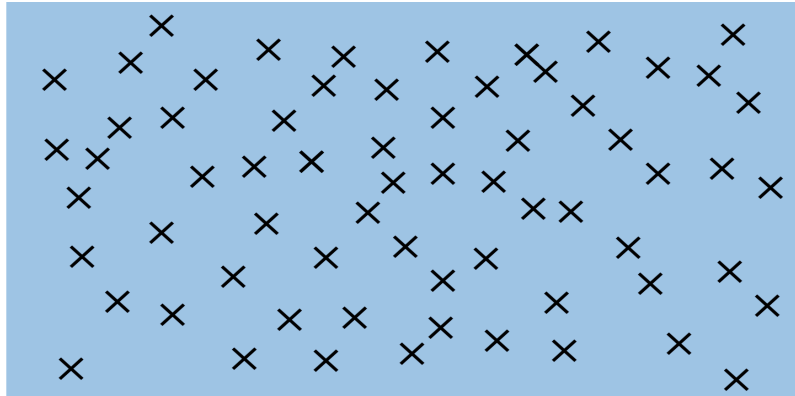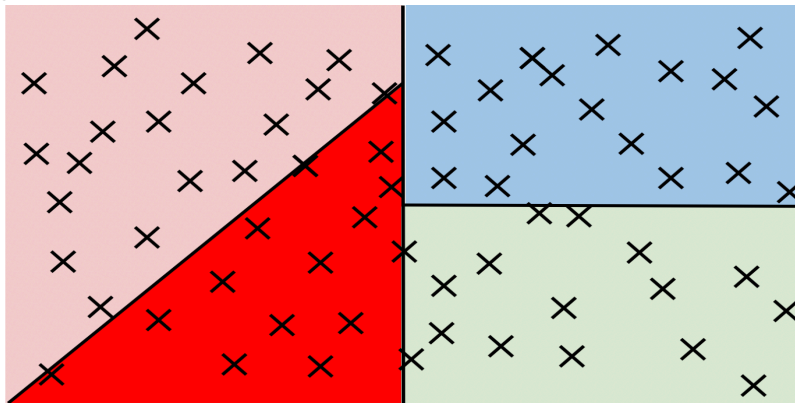Let's say your data points are plotted like this:

Figure 3

You can divide the vector space into regions and search within one region for nearest neighbors of a given vector.


```
In [39]:  N_VECS = len(all_tweets)      # This many vectors.
          N_DIMS = len(ind2Tweet[1])    # Vector dimensionality.
          print(f"Number of vectors is {N_VECS} and each has {N_DIMS} dimensions.")
```

Number of vectors is 10000 and each has 300 dimensions.

Choosing the number of planes

- Each plane divides the space to 2 parts.
- So n planes divide the space into 2^n hash buckets.
- We want to organize 10,000 document vectors into buckets so that every bucket has about 16 vectors.
- For that we need $\frac{10000}{16} = 625$ buckets.
- We're interested in n, number of planes, so that $2^n = 625$. Now, we can calculate $n = \log_2 625 = 9.29 \approx 10$.

```
In [40]:  # The number of planes. We use log2(625) to have ~16 vectors/bucket.
          N_PLANES = 10
          # Number of times to repeat the hashing to improve the search.
          N_UNIVERSES = 25
```

3.4 - Getting the Hash Number for a Vector

For each vector, we need to get a unique number associated to that vector in order to assign it to a "hash bucket".

Hyperplanes in Vector Spaces

- In 3-dimensional vector space, the hyperplane is a regular plane. In 2 dimensional vector space, the hyperplane is a line.
- Generally, the hyperplane is subspace which has dimension 1 lower than the original vector space has.
- A hyperplane is uniquely defined by its normal vector.
- Normal vector n of the plane π is the vector to which all vectors in the plane π are orthogonal (perpendicular in 3 dimensional case).

Using Hyperplanes to Split the Vector Space

We can use a hyperplane to split the vector space into 2 parts.

- All vectors whose dot product with a plane's normal vector is positive are on one side of the plane.
- All vectors whose dot product with the plane's normal vector is negative are on the other side of the plane.

Encoding Hash Buckets

- For a vector, we can take its dot product with all the planes, then encode this information to assign the vector to a single hash bucket.
- When the vector is pointing to the opposite side of the hyperplane than normal, encode it by 0.
- Otherwise, if the vector is on the same side as the normal vector, encode it by 1.
- If you calculate the dot product with each plane in the same order for every vector, you've encoded each vector's unique hash ID as a binary number, like [0, 1, 1, ..., 0]

Exercise 9 - hash_value_of_vector

We've initialized hash table `hashes` for you. It is list of `N_UNIVERSES` matrices, each describes its own hash table. Each matrix has `N_DIMS` rows and `N_PLANES` columns. Every column of that matrix is a `N_DIMS`-dimensional normal vector for each of `N_PLANES` hyperplanes which are used for creating buckets of the particular hash table.

Exercise: Your task is to complete the function `hash_value_of_vector` which places vector `v` in the correct hash bucket.

- First multiply your vector `v`, with a corresponding plane. This will give you a vector of dimension $(1, \text{N\_planes})$.
- You will then convert every element in that vector to 0 or 1.
- You create a hash vector by doing the following: if the element is negative, it becomes a 0, otherwise you change it to a 1.
- You then compute the unique number for the vector by iterating over `N_PLANES`
- Then you multiply 2^i times the corresponding bit (0 or 1).
- You will then store that sum in the variable `hash_value`.

Intructions: Create a hash for the vector in the function below. Use this formula:

$$hash = \sum_{i=0}^{N-1} \left(2^i \times h_i\right)$$

Create the sets of planes

- Create multiple (25) sets of planes (the planes that divide up the region).
- You can think of these as 25 separate ways of dividing up the vector space with a different set of planes.
- Each element of this list contains a matrix with 300 rows (the word vector have 300 dimensions), and 10 columns (there are 10 planes in each "universe").

```
In [41]: np.random.seed(0)
         planes_l = [np.random.normal(size=(N_DIMS, N_PLANES))
                     for _ in range(N_UNIVERSES)]
```

Hints

```
In [42]: # UNQ_C17
         import numpy as np

         def hash_value_of_vector(v, planes):
             """Create a hash for a vector; hash_id says which random hash to use.
             Input:
                 - v: vector of tweet, shape (1, N_DIMS)
                 - planes: matrix of shape (N_DIMS, N_PLANES)
             Output:
                 - hash_value: integer used as the hash for the vector
             """
             ### START CODE HERE ###
             dot_product = np.dot(v, planes)                    # shape (1, N_PLANES)
             sign_of_dot_product = np.sign(dot_product)         # -1, 0, 1
             h = (sign_of_dot_product >= 0).astype(int)         # convert to 0/1, tred
             h = h.flatten()                                    # convert (1, N_PLANES

             hash_value = 0
             n_planes = h.shape[0]
             for i in range(n_planes):
                 hash_value += (2 ** i) * h[i]
             ### END CODE HERE ###

             hash_value = int(hash_value)
             return hash_value
```

```
In [43]:  # UNQ_C18 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
          # You do not have to input any code in this cell, but it is relevant to gradin

          np.random.seed(0)
          idx = 0
          planes = planes_l[idx]   # get one 'universe' of planes to test the function
          vec = np.random.rand(1, 300)
          print(f" The hash value for this vector,",
                f"and the set of planes at index {idx},",
                f"is {hash_value_of_vector(vec, planes)}")
```

```
 The hash value for this vector, and the set of planes at index 0, is 768
```

Expected Output

```
The hash value for this vector, and the set of planes at index 0, is
768
```
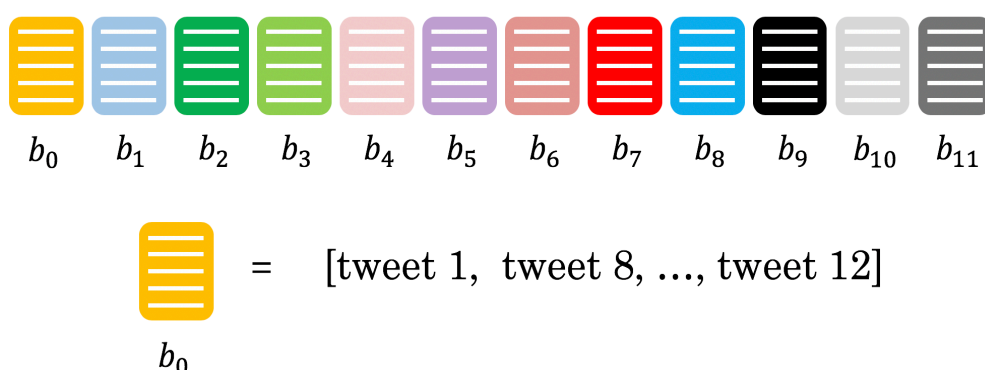
```
In [44]:  # Test your function
          w4_unittest.test_hash_value_of_vector(hash_value_of_vector)
```

```
All tests passed
```

3.5 - Creating a Hash Table

Exercise 10 - make_hash_table

Given that you have a unique number for each vector (or tweet), You now want to create a hash table. You need a hash table, so that given a hash_id, you can quickly look up the corresponding vectors. This allows you to reduce your search by a significant amount of time.

$b_0 \quad b_1 \quad b_2 \quad b_3 \quad b_4 \quad b_5 \quad b_6 \quad b_7 \quad b_8 \quad b_9 \quad b_{10} \quad b_{11}$

 $=$ [tweet 1, tweet 8, ..., tweet 12]

b_0

We have given you the `make_hash_table` function, which maps the tweet vectors to a bucket and stores the vector there. It returns the `hash_table` and the `id_table`. The `id_table` allows you know which vector in a certain bucket corresponds to what tweet.

Hints

```
In [45]:  # UNQ_C19 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)

          # This is the code used to create a hash table:
          # This function is already implemented for you. Feel free to read over it.

          ### YOU CANNOT EDIT THIS CELL

          def make_hash_table(vecs, planes, hash_value_of_vector=hash_value_of_vector):
              """
              Input:
                  - vecs: list of vectors to be hashed.
                  - planes: the matrix of planes in a single "universe", with shape (emb
              Output:
                  - hash_table: dictionary - keys are hashes, values are lists of vector
                  - id_table: dictionary - keys are hashes, values are list of vectors i
                                  (it's used to know which tweet corresponds to the
              """
              # number of planes is the number of columns in the planes matrix
              num_of_planes = planes.shape[1]

              # number of buckets is 2^(number of planes)
              # ALTERNATIVE SOLUTION COMMENT:
              # num_buckets = pow(2, num_of_planes)
              num_buckets = 2**num_of_planes

              # create the hash table as a dictionary.
              # Keys are integers (0,1,2.. number of buckets)
              # Values are empty lists
              hash_table = {i: [] for i in range(num_buckets)}

              # create the id table as a dictionary.
              # Keys are integers (0,1,2... number of buckets)
              # Values are empty lists
              id_table = {i: [] for i in range(num_buckets)}

              # for each vector in 'vecs'
              for i, v in enumerate(vecs):
                  # calculate the hash value for the vector
                  h = hash_value_of_vector(v, planes)

                  # store the vector into hash_table at key h,
                  # by appending the vector v to the list at key h
                  hash_table[h].append(v) # @REPLACE None

                  # store the vector's index 'i' (each document is given a unique intege
                  # the key is the h, and the 'i' is appended to the list at key h
                  id_table[h].append(i) # @REPLACE None

              return hash_table, id_table
```

```
In [46]:  # UNQ_C20 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
          # You do not have to input any code in this cell, but it is relevant to gradin
          planes = planes_l[0]   # get one 'universe' of planes to test the function
          tmp_hash_table, tmp_id_table = make_hash_table(document_vecs, planes)

          print(f"The hash table at key 0 has {len(tmp_hash_table[0])} document vectors"
          print(f"The id table at key 0 has {len(tmp_id_table[0])} document indices")
          print(f"The first 5 document indices stored at key 0 of id table are {tmp_id_t
```

```
The hash table at key 0 has 3 document vectors
The id table at key 0 has 3 document indices
The first 5 document indices stored at key 0 of id table are [3276, 3281, 32
82]
```

Expected output

```
The hash table at key 0 has 3 document vectors
The id table at key 0 has 3 document indices
The first 5 document indices stored at key 0 of id table are [3276, 3
281, 3282]
```

```
In [47]:  # Test your function
          w4_unittest.test_make_hash_table(make_hash_table)
```

```
All tests passed
```

3.6 - Creating all Hash Tables

You can now hash your vectors and store them in a hash table that would allow you to quickly look up and search for similar vectors. Run the cell below to create the hashes. By doing so, you end up having several tables which have all the vectors. Given a vector, you then identify the buckets in all the tables. You can then iterate over the buckets and consider much fewer vectors. The more tables you use, the more accurate your lookup will be, but also the longer it will take.

```
In [48]:  # Creating the hashtables
          def create_hash_id_tables(n_universes):
              hash_tables = []
              id_tables = []
              for universe_id in range(n_universes):  # there are 25 hashes
                  print('working on hash universe #:', universe_id)
                  planes = planes_l[universe_id]
                  hash_table, id_table = make_hash_table(document_vecs, planes)
                  hash_tables.append(hash_table)
                  id_tables.append(id_table)

              return hash_tables, id_tables

          hash_tables, id_tables = create_hash_id_tables(N_UNIVERSES)
```

```
working on hash universe #: 0
working on hash universe #: 1
working on hash universe #: 2
working on hash universe #: 3
working on hash universe #: 4
working on hash universe #: 5
working on hash universe #: 6
working on hash universe #: 7
working on hash universe #: 8
working on hash universe #: 9
working on hash universe #: 10
working on hash universe #: 11
working on hash universe #: 12
working on hash universe #: 13
working on hash universe #: 14
working on hash universe #: 15
working on hash universe #: 16
working on hash universe #: 17
working on hash universe #: 18
working on hash universe #: 19
working on hash universe #: 20
working on hash universe #: 21
working on hash universe #: 22
working on hash universe #: 23
working on hash universe #: 24
```

Approximate K-NN

Exercise 11 - approximate_knn

Implement approximate K nearest neighbors using locality sensitive hashing, to search for documents that are similar to a given document at the index `doc_id` .

Inputs

- `doc_id` is the index into the document list `all_tweets` .
- `v` is the document vector for the tweet in `all_tweets` at index `doc_id` .
- `planes_l` is the list of planes (the global variable created earlier).
- `k` is the number of nearest neighbors to search for.
- `num_universes_to_use` : to save time, we can use fewer than the total number of available universes. By default, it's set to `N_UNIVERSES` , which is 25 for this assignment.

- `hash_tables` : list with hash tables for each universe.
- `id_tables` : list with id tables for each universe.

The `approximate_knn` function finds a subset of candidate vectors that are in the same "hash bucket" as the input vector 'v'. Then it performs the usual k-nearest neighbors search on this subset (instead of searching through all 10,000 tweets).

Hints

In [49]:
```python
# UNQ_C21
def approximate_knn(doc_id, v, planes_l, hash_tables, id_tables, k=1, num_univ
    vecs_to_consider_l = list()
    ids_to_consider_l = list()
    ids_to_consider_set = set()

    for universe_id in range(num_universes_to_use):
        planes = planes_l[universe_id]
        hash_value = hash_value_of_vector(v, planes)
        hash_table = hash_tables[universe_id]
        document_vectors_l = hash_table[hash_value]
        id_table = id_tables[universe_id]
        new_ids_to_consider = id_table[hash_value]

        ### START CODE HERE ###
        for i, new_id in enumerate(new_ids_to_consider):
            if doc_id == new_id:
                continue
            if new_id not in ids_to_consider_set:
                document_vector_at_i = document_vectors_l[i]
                vecs_to_consider_l.append(document_vector_at_i)
                ids_to_consider_l.append(new_id)
                ids_to_consider_set.add(new_id)
        ### END CODE HERE ###

    print("Fast considering %d vecs" % len(vecs_to_consider_l))
    vecs_to_consider_arr = np.array(vecs_to_consider_l)
    nearest_neighbor_idx_l = nearest_neighbor(v, vecs_to_consider_arr, k=k)
    nearest_neighbor_ids = [ids_to_consider_l[idx] for idx in nearest_neighbor

    return nearest_neighbor_ids
```

In [50]:
```python
#document_vecs, ind2Tweet
doc_id = 0
doc_to_search = all_tweets[doc_id]
vec_to_search = document_vecs[doc_id]
```

In [51]:
```python
# UNQ_C22 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# You do not have to input any code in this cell, but it is relevant to gradin

# Sample
nearest_neighbor_ids = approximate_knn(
    doc_id, vec_to_search, planes_l, hash_tables, id_tables, k=3, num_universe
```

Fast considering 77 vecs

```
In [52]: print(f"Nearest neighbors for document {doc_id}")
         print(f"Document contents: {doc_to_search}")
         print("")

         for neighbor_id in nearest_neighbor_ids:
             print(f"Nearest neighbor at document id {neighbor_id}")
             print(f"document contents: {all_tweets[neighbor_id]}")
```

Nearest neighbors for document 0
Document contents: #FollowFriday @France_Inte @PKuchly57 @Milipol_Paris for
being top engaged members in my community this week :)

Nearest neighbor at document id 51
document contents: #FollowFriday @France_Espana @reglisse_menthe @CCI_inter
for being top engaged members in my community this week :)
Nearest neighbor at document id 2478
document contents: #ShareTheLove @oymgroup @musicartisthere for being top Hi
ghValue members this week :) @nataliavas http://t.co/IWSDMtcayt (http://t.c
o/IWSDMtcayt)
Nearest neighbor at document id 105
document contents: #FollowFriday @straz_das @DCarsonCPA @GH813600 for being
top engaged members in my community this week :)

```
In [53]: # Test your function
         w4_unittest.test_approximate_knn(approximate_knn, hash_tables, id_tables)
```

Fast considering 77 vecs
Fast considering 153 vecs
 All tests passed

4 Conclusion

Congratulations - Now you can look up vectors that are similar to the encoding of your tweet
using LSH!